

Week 8 Code

Ryan Martin

February 25, 2018

Notes for Week 8

Exam Chapters 4,5 and 6 (Homework 3,4 and 5)

- If performing an F test for one variable, you can check your answer with t-test from the unrestricted regression.
- Jarque-Bera is, under the null that the errors are normal, distributed according to a chi-squared distribution with 2 degrees of freedom ($\chi^2(2)$). The two degrees of freedom is because the statistic is the sum of two squared normals under the null (S^2 and $(K-3)^2$). (You may recall from econ 41 that a chi-squared distribution with n degrees of freedom can be constructed as the sum of n independent normal random variables). Rejecting the null means, given the data, the errors are unlikely to be normally distributed.
- If \hat{y}_n denotes the natural predictor, we can define the generalized R^2 for the log-linear model as $R_g^2 := r_{y\hat{y}_n}^2$. This is in contrast to the standard R^2 , mentioned above, which would be $R^2 = r_{\log y, \widehat{\log(y)}}^2$. Further note that

$$\text{corr}(aX, Y) = \frac{\text{cov}(aX, Y)}{\sqrt{\text{var}(aX)\text{var}(Y)}} = \frac{a\text{cov}(X, Y)}{|a|\sqrt{\text{var}(X)\text{var}(Y)}} = \text{sign}(a)\text{corr}(X, Y)$$

where

$$\text{sign}(a) = \begin{cases} 1 & \text{if } a > 0 \\ 0 & \text{if } a = 0 \\ -1 & \text{if } a < 0 \end{cases}$$

Finally, we can therefore conclude, that since $e^{\hat{\sigma}^2/2}y_n = y_c$ and $e^{\hat{\sigma}^2/2} > 0$, that $\text{corr}(\hat{y}_n, y) = \text{corr}(\hat{y}_c, y)$ and so $R_g^2 = r_{y, \hat{y}_n}^2 = r_{y, \hat{y}_c}^2$ for the linear and log-linear models.

Model Selection:

- Higher R^2 because larger means explaining more variation in outcome variable
- Higher adjusted R^2 (denoted \bar{R}^2), which tries to account for advantage larger models have. (will always improve if t-stat from added variable is larger in absolute value than 1)
- $\bar{R}^2 = 1 - \frac{SSE \times (N-1)}{SST \times (N-K)}$ adjusted R squared. Bigger is better. It rescales the term being subtracted by $\frac{N-1}{N-K}$. Note, since $N-1 > N-K$, subtracting a bigger number, so $R^2 \geq \bar{R}^2$. The more parameters in your model, the larger the fraction $\frac{N-1}{N-K}$ but also the lower your SSE . Fact: \bar{R}^2 will increase with one more coefficient if that coefficients t-value is greater than 1 (vs 2 for two-sided .95 confidence when N large). Therefore, interpreting it as model selection not always best idea.

AIC and BIC/SC - smaller (or more negative) is better (for both).

RESET - include \hat{y}^k for $k = 2, 3, \dots$ in your model. Test the regression coefficients of your \hat{y}^k using F-test vs model without \hat{y}^k term. If reject F-test null (that all \hat{y}^k coefficients are zero), then your model is “not adequate”.

Collinearity

- Note, having collinearity isn't necessarily a problem. It's only a problem because you will get poorer variance estimates of coefficients.
- When you have (non-perfect) collinearity, inclusion or exclusion of some regressors will strongly influence coefficient estimates.
- To test for collinearity between two variables x and z , you can look at their correlation.
- To test for collinearity between three or more (explanatory/independent) variables, x_1, x_2, x_3 , etc., regress x_1 on x_2 and x_3 and then x_2 on x_1 and x_3 , etc. R^2 in these regressions above .80 suggest collinearity. This is called "auxiliary regressions"
- An alternative way to report these auxiliary regression is with VIF (Variance Inflation Factors). See a VIF example in 6.22 part (f) from last week's lecture notes, or from your class notes. To understand the relationship between VIF and auxiliary regressions, consider the auxiliary regression of x_1 on x_2, x_3, \dots, x_K where x_K is the last regressor. Let R_1^2 be the R squared from this regression. Then $VIF_1 = \frac{1}{1-R_1^2}$. We can calculate this for each auxiliary regression 1 to K . A high R^2 is greater than .8. A high VIF comes from an R^2 greater than .8. So a high VIF would be bigger than or equal to $\frac{1}{1-.8} = 5$

R^2

with two term regression (that is, with just a constant and coefficient on scalar x explaining y), $R^2 = r_{xy}^2 = \frac{s_{xy}}{s_x s_y}$ where r_{xy} is the estimated correlation between x and y , which is the estimated covariance of x and y divided by the standard errors. With more terms, $R^2 = r_{y\hat{y}}$

Problems from Changing Scale of x, y

You can infer all the answers from the formulas on the formula sheet:

$$b_2 = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{\sum (x_i - \bar{x})^2} \text{ and } b_1 = \bar{y} - b_2 \bar{x} \text{ or just memorize the rules from the book page 140.}$$

For example, if x is replaced by $x^* = 5x$, then trivially $\bar{x}^* = 5\bar{x}$ and b_2 turns into

$$b_2^* = \frac{\sum (5x_i - 5\bar{x})(y_i - \bar{y})}{\sum (5x_i - 5\bar{x})^2} = \frac{5 \sum (x_i - \bar{x})(y_i - \bar{y})}{25 \sum (x_i - \bar{x})^2} = \frac{1}{5} b_2$$

and

$$b_1 = \bar{y} - b_2^* \bar{x}^* = \bar{y} - \frac{1}{5} b_2 (5\bar{x}) = \bar{y} - b_2 \bar{x}$$

That is, the new b_2 is $1/5$ the old b_2 and the new b_1 is the same as the old b_1 .

Similarly, if multiply the y 's by 5 such that $y^* = 5y$, without changing the x 's, then

$$b_2^* = \frac{\sum (x_i - \bar{x})(5y_i - 5\bar{y})}{\sum (x_i - \bar{x})^2} = 5b_2$$

and

$$b_1 = \bar{y}^* - b_2^* \bar{x} = 5\bar{y} - 5b_2 \bar{x} = 5(\bar{y} - b_2 \bar{x})$$

That is, the new b_1 and b_2 are both just rescaled by 5.

You can plug these formulas into the formulas for t statistics and R^2 values to see these won't change under rescalings. And since t statistics aren't changing, p values don't change either.

Source	SS	df	MS	Number of obs = 506		
Model	2.4135e+10	3	8.0450e+09	F(3, 502) = 216.08		
Residual	1.8691e+10	502	37232218.6	Prob > F = 0.0000		
				R-squared = 0.5636		
				Adj R-squared = 0.5610		
				Root MSE = 6101.8		
Total	4.2826e+10	505	84803032			

price	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]	
lnox	-7579.001	1537.677	-4.93	0.000	-10600.08	-4557.926
rooms	7928.224	407.9288	19.44	0.000	7126.766	8729.682
crime	-197.617	35.18799	-5.62	0.000	-266.7508	-128.4831
_cons	-13764.21	4054.979	-3.39	0.001	-21731.04	-5797.393

Figure 1:

Equations can you get from output table?

- $SST = SSE + SSR$
- $SST = \sum (y_i - \bar{y})^2$ (Total Variation)
- $SSR = \sum (\hat{y}_i - \bar{y})^2$ Sum of Squared Regression (explained variation)
- $SSE = \sum \hat{e}_i^2 = \sum (y - \hat{y}_i)^2$ (Unexplained Variation)

Note, in this table: - $SSR = 2.41 \times 10^{10}$ - $SST = 4.282 \times 10^{10}$ - $SSE = 1.86 \times 10^{10}$

- $F(3, 502)$ refers to the f statistic of the model versus a constant. Prob > F refers to the p value of this F statistic. A low p-value (below .05) suggests the model does a statistically significantly better job at explaining the model's variation than the job a constant would. (i.e., a large p-value says your model is garbage, you can't even beat a constant's predictive value in a statistically significant way)
- $Root\ MSE = \sqrt{\frac{1}{N-K} SSE} = \hat{\sigma}$
- Of course, N is the number of observations (upper right) and counting the coefficients in the regression gets you K
- Note that the t-statistics tells you the answer to F-statistics where restrict just that single coefficient. For example, the F statistic for the hypothesis of removing rooms would be $(19.44)^2$ and it's p-value would be 0.000. This only works when testing to remove 1 term.
- There are also the coefficient, standard error, t-value, p-value and confidence interval relationships learned for the first exam in this table.

Omitted Variable Bias

- If $y = \beta_1 + \beta_2 x_2 + \beta_3 x_3 + e$, but you don't include x_3 in the regression, and just regress to get $y = b_1 + b_2 x_2$, then $bias(b_2) = E(b_2) - \beta_2 = \beta_3 \frac{cov(x_2, x_3)}{var(x_2)} = \beta_3 \frac{r_{x_2, x_3}}{\sigma_2^2}$. So the sign of the bias depends on the (unknown but estimatable if you have x_3) true coefficient β_3 and the sign of the sample correlation of x_2 and x_3 . Variance is always positive so doesn't affect the sign

New Code

```
nlcom (vs lincom)
  - performs delta-method test on nonlinear combo
  of variables
test
  - performs an F test on the listed variables.
estat vce
  - to get covariance matrix of recent regressors
corr
  - to get correlation matrix of recent regressors
  - can specify only certain values, like correlate wage yhatn yhatc
  - on page 150

reg y c.x1## c.x2
  - ## is the same as reg y x1 x2 c.x1#c.x2

ttail(df, t1)
  - prob of t random var with degrees of freedom df is larger than t1

invttail(df, 05)
  - finds quantile t1 such that prob to the *right* of t1 is .05

chi2tail(2, jb) //page 137
  - to get pvalue to the right of chi squared at jb
invchi2tail(2, 05) //page 136
  - to get cutoff value jb such that .05 to the right of chi squared
  - of degrees of freedom 2.

Ftail(J, N-k, fstat)

invFtail(J,N_K, alpha)

cnsreg
  - regression where add (linear) constraints

predict yhat
  - predicts dependent variable at all x values in table, calls
  - that predicted vector yhat
  - same as predict yhat, xb (xb is default option)
predict ehat, residuals
  - calculates the residuals at all x values in table, calls ehat
predict expectvalse, stdp
  - calculates standard error of predicted values at all points
  - this is variance of, for example se(b_1 + 3 b_2) if
  - x = 3 is in our table
  - think of as standard error for average
  - stores standard errors in expectvalse (can use different name)
predict sef, stdf
  - predicts standard forecast errors at all x values in table
  - stores these predicted variables in variable called sef
  - (note, each point has its own forecast standard error,
  - giving N forecast standard errors (see the formula has x in it)
```

```

- vs each coefficient estimate (e.g b_2) has only 1 standard error)
- standard forecast error larger than prediction standard errors!
- can see this in formula
estat ovtest
- does RESET test, including yhat^2, yhat^3 and yhat^4
  to the regression
r(skewness) //page 136
- r() gets results from general commands
- vs e() gets results from estimation commands
- vs s() from parsing commands, see later

// to see all possible, type
return list
ereturn list
sreturn list
// if see in code e(blah) or r(blah), just pulling a stored value

e(rss) - sum of squared residuals (SSE not SSR!)
e(rmse) - root mean squared error (Root MSE)
e(r2) - R squared value
e(r2_a) - adjusted R squared

_se[varname] //to get se of variable named varname
_b[varname] // to get estimate of variable named varname

margins, dydx(advert) at(atvert = .5 2)
- to estimate and CI for how y changes with respect to advert
- at .5 and at 2
- if no "at", then does at average advert, average all other vars
- if advert replaced with *, calculates for all dependent vars
- vs margins, eyex(advert) from chapter 2, computes elasticities

rvpplot
- residual vs predictor plot
- many

tway lfit
- calculates prediction for yvar from xvar, plots fitted line

i. #to specify a variable is an indicator
c. #to specify a variable is to be created as continuous

## example
regress mpg cyl eng wgt // p.204
test cyl eng //does joint test that cyl = 0 and eng = 0

##
reg sales price advert a2 //page 180
nlcom (1 - _b[advert])/(2*_b[a2])
// tests if 1 minus coefficient of advert divided by twice
// the coefficient of a2 is different from 0

```

```
//vs
lncom -.4*price + .8*advert, level(90) //page 179
//tests if price coefficient = 2*advertising coefficient or
//equivalently
//if -.4*price coefficient + .8 times advert coefficient
// is different from 0 at (two-sided) alpha = .10 level
```

- $F_{1-\alpha,1,n-k} = t_{1-\alpha/2,n-k}^2 = t_{\alpha/2,n-k}^2$ for all $n - k > 0$ F and t-tests give the same p-value and there is a relationship between the variables

```
#Comparing Problems
```

```
qf(.95, 1, 218)
```

```
## [1] 3.884469
```

```
qt(.975, 218)
```

```
## [1] 1.970906
```

```
qt(.975, 218)^2
```

```
## [1] 3.884469
```

```
qt(.025, 218)^2
```

```
## [1] 3.884469
```

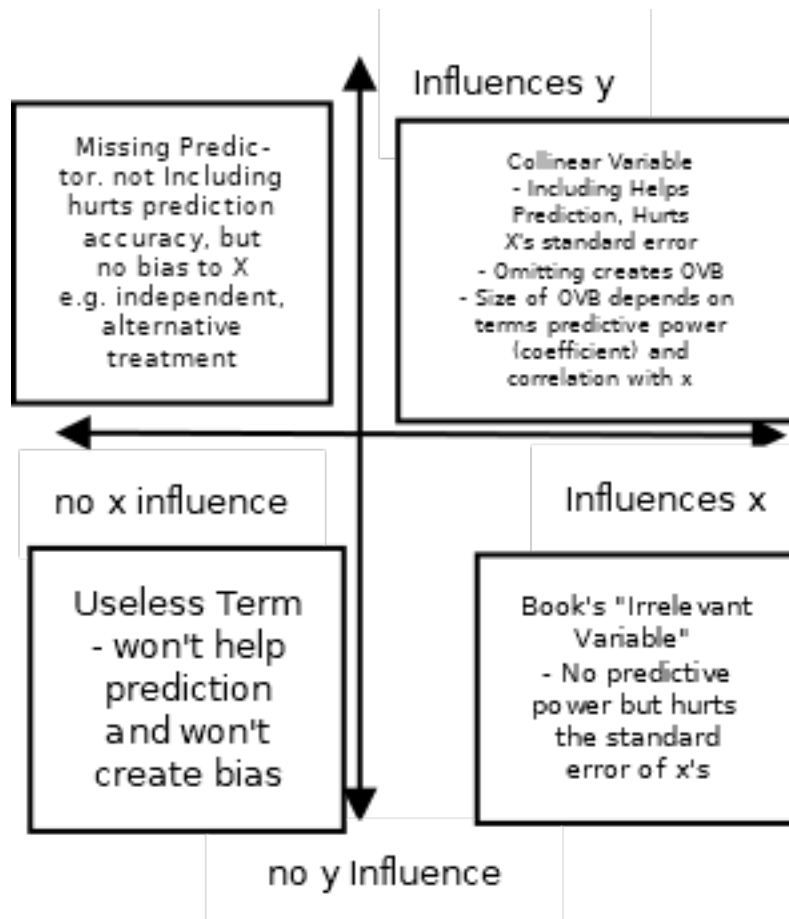


Figure 2: Summary Table